

# Exercises

## Debugging

Use `journalctl` on Linux or `log show` on macOS to get the super user accesses and commands in the last day. If there aren't any you can execute some harmless commands such as `sudo ls` and check again.

```
~$ man journalctl
# Show superuser accesses and commands in the last day
~$ journalctl -r --since yesterday --output=short | grep sudo
May 06 14:10:05 ramzel-Inspiron-14-3467 sudo[27380]: ramzel :
TTY=pts/0 ; PWD=/home/ramzel ; USER=root ; COMMAND=/usr/bin/apt
autoremove
May 06 14:09:55 ramzel-Inspiron-14-3467 sudo[27376]: ramzel :
TTY=pts/0 ; PWD=/home/ramzel ; USER=root ; COMMAND=/usr/bin/apt clean
May 06 14:09:23 ramzel-Inspiron-14-3467 sudo[27354]: ramzel :
TTY=pts/0 ; PWD=/home/ramzel ; USER=root ; COMMAND=/usr/bin/lshw -c
memory
May 06 14:06:09 ramzel-Inspiron-14-3467 sudo[27200]: ramzel :
TTY=pts/0 ; PWD=/home/ramzel ; USER=root ; COMMAND=/usr/bin/vim
/etc/samba/smb.conf
May 06 13:55:46 ramzel-Inspiron-14-3467 sudo[16514]: ramzel :
TTY=pts/0 ; PWD=/home/ramzel ; USER=root ; COMMAND=/usr/bin/apt
update
```

*I mostly used sudo to do apt update and to install some applications in the past day.*

Install [shellcheck](#) and try checking the following script. What is wrong with the code? Fix it. Install a linter plugin in your editor so you can get your warnings automatically.

```
#!/bin/sh
## Example: a typical script with several problems
for f in $(ls *.m3u)
# Iterating over ls output is fragile. Use globs.shellcheck(SC2045)

do
    grep -qi hq.*mp3 $f \
    # Quote the grep pattern so the shell won't interpret
    it.shellcheck(SC2062)
    # Double quote to prevent globbing and word splitting.shellcheck(SC2086)

    && echo -e 'Playlist $f contains a HQ file in mp3 format'
    # Expressions don't expand in single quotes, use double quotes for
    that.shellcheck(SC2016)
    # In POSIX sh, echo flags are undefined.shellcheck(SC3037)
done
```

*I used the ShellCheck extension for VSCode and it's a really awesome tool. While coding, it immediately flags programming errors and other problems with the code. ShellCheck tells you what is wrong with it and how to fix it.*

*Here's the working script:*

```
#!/bin/sh
## Fixed the script with several problems
for f in *.m3u
do
    [ -e "$f" ] || break
    grep -qi "hq.*mp3" "$f" \
    && echo "Playlist $f contains a HQ file in mp3 format"
done
```

*I ran it through the terminal after creating some .m3u files.*

```
$ chmod +x problem-script.sh
$ ./problem-script.sh
Playlist playlist-b.m3u contains a HQ file in mp3 format
Playlist playlist.m3u contains a HQ file in mp3 format
```

# Profiling

[Here](#) are some sorting algorithm implementations. Use [cProfile](#) and [line\\_profiler](#) to compare the runtime of insertion sort and quicksort. What is the bottleneck of each algorithm? Use then [memory\\_profiler](#) to check the memory consumption, why is insertion sort better? Check now the inplace version of quicksort. Challenge: Use [perf](#) to look at the cycle counts and cache hits and misses of each algorithm.

```
# Using cProfile to profile the runtime of the sorting functions
$ python -m cProfile -s tottime sorts.py
      398688 function calls (332330 primitive calls) in 0.243
seconds

      Ordered by: internal time

      ncalls  tottime  percall  cumtime  percall
filename:lineno(function)
      78438    0.065    0.000    0.070    0.000
random.py:177(randrange)
# quicksort function took 39 milliseconds
34064/1000    0.039    0.000    0.041    0.000 sorts.py:23(quicksort)
# quicksort_inplace ran for 27 milliseconds
34294/1000    0.027    0.000    0.031    0.000
sorts.py:32(quicksort_inplace)
78438    0.021    0.000    0.091    0.000 random.py:240(randint)
# insertionsort took 21 milliseconds
      1000    0.021    0.000    0.021    0.000
sorts.py:11(insertionsort)
...
```

*cProfile showed that the quick sort function took about 1.8 times longer to run than the insertion sort.*

*Taking a look at the runtime of the quicksort\_inplace function, it is a bit faster than quick sort but insertion sort still proved to be better.*

Using line\_profiler:

```
~$ pip install line_profiler
```

Add decorator above functions to profile

*insertionsort(array)*

```
@profile
def insertionsort(array):

    for i in range(len(array)):
        ...
```

*quicksort(array)*

```
@profile
def quicksort(array):
    if len(array) <= 1:
        return array
    ...
```

Running line\_profiler

```
# Profile functions quick sort and insertion sort
$ kernprof -l -v sorts.py
Wrote profile results to sorts.py.lprof
Timer unit: 1e-06 s
```

```
# Runtime of insertion sort
Total time: 0.219251 s
File: sorts.py
Function: insertionsort at line 11
```

Line #	Hits	Time	Per Hit	% Time	Line Contents
11					@profile
12					def
					insertionsort(array):
13					
14	26103	6922.0	0.3	3.1	for i in

```

range(len(array)):
    15      25103      6599.0      0.3      3.0      j = i-1
    16      25103      6893.0      0.3      3.1      v = array[i]
    17      228923     73297.0      0.3     33.3      while j >= 0
and v < array[j]:
    18      203820     62416.0      0.3     28.4
array[j+1] = array[j]
    19      203820     56112.0      0.3     25.5      j -= 1
    20      25103      7506.0      0.3      3.4      array[j+1] =
v
    21      1000       232.0      0.2      0.1      return array

```

```

# Runtime of quick sort
Total time: 0.099628 s
File: sorts.py
Function: quicksort at line 23

```

Line #	Hits	Time	Per Hit	% Time	Line Contents
24					@profile
25					def
quicksort(array):					
26	33716	15667.0	0.5	16.2	if len(array) <=
1:					
27	17358	6727.0	0.4	7.0	return array
28	16358	6797.0	0.4	7.0	pivot = array[0]
29	16358	26267.0	1.6	27.2	left = [i for i
					in array[1:] if i < pivot]
30	16358	25663.0	1.6	26.5	right = [i for i
					in array[1:] if i >= pivot]
31	16358	15602.0	1.0	16.1	return
quicksort(left) + [pivot] + quicksort(right)					

## Memory Profile

Repeated memory profiling of functions quick sort and insertion sort consistently showed that insertion sort used less memory than insertion sort. However, the difference in memory usage isn't that significant at a thousand iterations.

Quick sort	39.44 MiB	39.48 MiB	39.41 MiB
Insertion sort	39.07 MiB	39.285 MiB	39.289 MiB

```
~$ pip install -U memory_profiler
```

```
# Memory profile of quicksort
```

```
$ python3.6 -m memory_profiler sorts.py
```

```
Filename: sorts.py
```

Line #	Mem usage	Increment	Occurences	Line Contents
24	39.414 MiB	39.273 MiB	34758	@profile
25				def quicksort(array):
26	39.414 MiB	0.141 MiB	34758	if len(array) <=
1:				
27	39.414 MiB	0.000 MiB	17879	return array
28	39.414 MiB	0.000 MiB	16879	pivot = array[0]
29	39.414 MiB	0.000 MiB	162711	left = [i for i in
				array[1:] if i < pivot]
30	39.414 MiB	0.000 MiB	162711	right = [i for i
				in array[1:] if i >= pivot]
31	39.414 MiB	0.000 MiB	16879	return
				quicksort(left) + [pivot] + quicksort(right)

```
# Memory profile of insertionsort
```

```
$ python3.6 -m memory_profiler sorts.py
```

```
Filename: sorts.py
```

Line #	Mem usage	Increment	Occurences	Line Contents
12	39.070 MiB	39.070 MiB	1000	@profile
13				def
				insertionsort(array):
14				
15	39.070 MiB	0.000 MiB	27026	for i in
				range(len(array)):
16	39.070 MiB	0.000 MiB	26026	j = i-1
17	39.070 MiB	0.000 MiB	26026	v = array[i]
18	39.070 MiB	0.000 MiB	243316	while j >= 0
				and v < array[j]:
19	39.070 MiB	0.000 MiB	217290	array[j+1]
				= array[j]
20	39.070 MiB	0.000 MiB	217290	j -= 1

```
21    39.070 MiB    0.000 MiB    26026    array[j+1] = v
22    39.070 MiB    0.000 MiB    1000    return array
```

A common issue is that a port you want to listen on is already taken by another process. Let's learn how to discover that process pid. First execute `python -m http.server 4444` to start a minimal web server listening on port 4444. On a separate terminal run `lsof | grep LISTEN` to print all listening processes and ports. Find that process pid and terminate it by running `kill <PID>`.

```
# Start a web server
$ python3 -m http.server 4444
Serving HTTP on 0.0.0.0 port 4444 (http://0.0.0.0:4444/) ...
# Find its process id
$ sudo lsof | grep ":4444 .LISTEN"
python3    3934                ramzel    3u        IPv4
1150791    0t0          TCP *:4444 (LISTEN)
# Terminate the process
~$ kill 3934
```

```
$ python3 -m http.server 4444
Serving HTTP on 0.0.0.0 port 4444 (http://0.0.0.0:4444/) ...
Terminated
```